

```
In [1]: import numpy as np
import nnfs
from nnfs.datasets import spiral_data
import matplotlib.pyplot as plt
```

```
In [2]: import numpy as np
import cv2
import os

# Loads a MNIST dataset
def load_mnist_dataset(dataset, path):
    """
    returns samples and lables specified by path and dataset params

    loop through each label and append image to X and label to y
    """
    labels = os.listdir(os.path.join(path, dataset))

    X = []
    y = []

    for label in labels:
        image_counter = 0
        for file in os.listdir(os.path.join(path, dataset, label)):
            image = cv2.imread(os.path.join(path, dataset, label, file), cv2.IMREAD_UNCHANGED)
            X.append(image)
            y.append(label)

    return np.array(X), np.array(y).astype('uint8')

def create_data_mnist(path):
    """
    returns train X, y and test X and y
    """
    X, y = load_mnist_dataset('train', path)
    X_test, y_test = load_mnist_dataset('test', path)

    return X, y, X_test, y_test
```

```
In [3]: X, y, X_test, y_test = create_data_mnist('fashion_mnist_images')

# Scale features
X = (X.astype(np.float32) - 127.5) / 127.5
X_test = (X_test.astype(np.float32) - 127.5) / 127.5

# print(X.min(), X.max())
# print(X.shape)
```

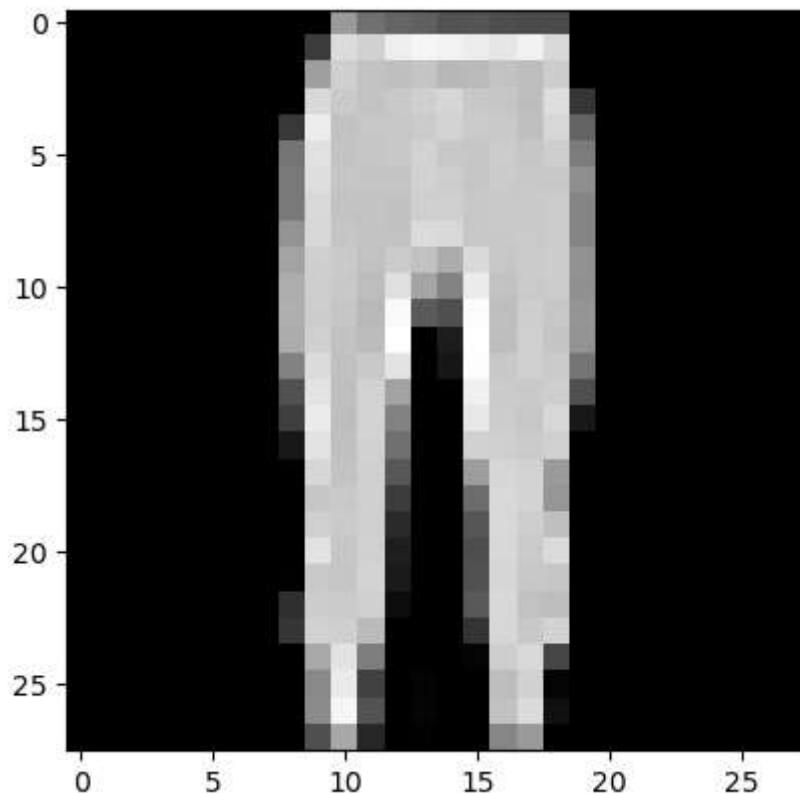
```
In [4]: print(X.shape)
print(y.shape)

(60000, 28, 28)
(60000,)
```

```
In [5]: # Reshape to vectors
X = X.reshape(X.shape[0], -1)
X_test = X_test.reshape(X_test.shape[0], -1)
```

```
In [15]: keys = np.array(range(X.shape[0]))
print(keys[:10])
np.random.shuffle(keys)
X = X[keys]
y = y[keys]
plt.imshow((X[4].reshape(28, 28)), cmap='gray')
plt.show()
```

[0 1 2 3 4 5 6 7 8 9]



```
In [7]: class Layer_Dense:
    def __init__(self, n_inputs, n_neurons,
                 weight_regularizer_l1=0, weight_regularizer_l2=0,
                 bias_regularizer_l1=0, bias_regularizer_l2=0):
        """
        n_inputs represents the no of features
        n_neurons represents the no of neurons we want in this particular layer

        weights are inputs x no of neurons - helps avoid transpose operation in dot pr
        weights range between -1 to +1 so tht they are close to each other and NN does

        biases is set to zero initially and if we encounter error where the entore out
        NN is zero we can intialize it to some value to avoid dead ANN

        rest four parameters refers to lambda that will be used for
        L1 and L2 regularization
        """
        self.weights = 0.01 * np.random.randn(n_inputs, n_neurons)
        self.biases = np.zeros((1, n_neurons))

        # Set regularization strength
        self.weight_regularizer_l1 = weight_regularizer_l1
        self.weight_regularizer_l2 = weight_regularizer_l2
        self.bias_regularizer_l1 = bias_regularizer_l1
```

```

self.bias_regularizer_l2 = bias_regularizer_l2

def forward(self, inputs):
    self.inputs = inputs
    self.output = np.dot(inputs, self.weights) + self.biases

def backward(self, dvalues):
    # Gradients on parameters
    self.dweights = np.dot(self.inputs.T, dvalues)
    self.dbiases = np.sum(dvalues, axis=0, keepdims=True)

    # Gradients on regularization
    # L1 on weights
    if self.weight_regularizer_l1 > 0:
        dL1 = np.ones_like(self.weights)
        dL1[self.weights < 0] = -1
        self.dweights += self.weight_regularizer_l1 * dL1

    # L2 on weights
    if self.weight_regularizer_l2 > 0:
        self.dweights += 2 * self.weight_regularizer_l2 * \
            self.weights

    # L1 on biases
    if self.bias_regularizer_l1 > 0:
        dL1 = np.ones_like(self.biases)
        dL1[self.biases < 0] = -1
        self.dbiases += self.bias_regularizer_l1 * dL1

    # L2 on biases
    if self.bias_regularizer_l2 > 0:
        self.dbiases += 2 * self.bias_regularizer_l2 * \
            self.biases

    # Gradient on values
    self.dinputs = np.dot(dvalues, self.weights.T)

```

In [8]:

```

class Activation_ReLU:
    def forward(self, inputs):
        self.inputs = inputs
        self.output = np.maximum(0, inputs)

    def backward(self, dvalues):
        # Since we need to modify original variable,
        # let's make a copy of values first
        self.dinputs = dvalues.copy()
        # Zero gradient where input values were negative
        self.dinputs[self.inputs <= 0] = 0

class Activation_Softmax:
    def forward(self, inputs):
        # Remember input values
        self.inputs = inputs

        exp_values = np.exp(inputs - np.max(inputs, axis=1, keepdims=True))
        probabilities = exp_values / np.sum(exp_values, axis=1, keepdims=True)
        self.output = probabilities

    # Backward pass

```

```

def backward(self, dvalues):
    # Create uninitialized array
    self.dinputs = np.empty_like(dvalues)

    # Enumerate outputs and gradients
    for index, (single_output, single_dvalues) in enumerate(zip(self.output, dvalues)):
        # Flatten output array
        single_output = single_output.reshape(-1, 1)
        # Calculate Jacobian matrix of the output and
        jacobian_matrix = np.diagflat(single_output) - np.dot(single_output, single_output.T)
        # Calculate sample-wise gradient
        # and add it to the array of sample gradients
        self.dinputs[index] = np.dot(jacobian_matrix, single_dvalues)

```

In [9]:

```

class Loss:
    def calculate(self, output, y):
        sample_losses = self.forward(output, y)
        mean_loss = np.mean(sample_losses)
        return mean_loss

    def regularization_loss(self, layer):

        regularization_loss = 0

        # L1 regularization - weights
        # calculate only when factor greater than 0
        if layer.weight_regularizer_l1 > 0:
            regularization_loss += layer.weight_regularizer_l1 * np.sum(np.abs(layer.weights))

        # L2 regularization - weights
        if layer.weight_regularizer_l2 > 0:
            regularization_loss += layer.weight_regularizer_l2 * np.sum(layer.weights * layer.weights)

        # L1 regularization - biases
        # calculate only when factor greater than 0
        if layer.bias_regularizer_l1 > 0:
            regularization_loss += layer.bias_regularizer_l1 * np.sum(np.abs(layer.biases))

        # L2 regularization - biases
        if layer.bias_regularizer_l2 > 0:
            regularization_loss += layer.bias_regularizer_l2 * np.sum(layer.biases * layer.biases)

        return regularization_loss

class Loss_CategoricalCrossentropy(Loss):
    def forward(self, y_pred, y_true):
        samples = len(y_pred)

        # account for zero values -log(0) = inf and then remove bias caused by its infinity
        y_pred_clipped = np.clip(y_pred, 1e-7, 1-1e-7)

        if len(y_true.shape) == 1:
            # y_true is in the form of [1, 0, 1, 1 ....]
            correct_confidences = y_pred_clipped[range(samples), y_true]

        elif len(y_true.shape) == 2:
            # y_true is in the form of matrix [[0, 1, 0], [1, 0, 0], [0, 1, 0], [0, 1, 0]]
            correct_confidences = np.sum(y_pred_clipped * y_true, axis=1)

```

```

negative_log_likelihood = - np.log(correct_confidences)

return negative_log_likelihood

def backward(self, dvalues, y_true):
    # Number of samples
    samples = len(dvalues)
    # Number of Labels in every sample
    # We'll use the first sample to count them
    labels = len(dvalues[0])
    # If Labels are sparse, turn them into one-hot vector
    if len(y_true.shape) == 1:
        y_true = np.eye(labels)[y_true]
    # Calculate gradient
    self.dinputs = -y_true / dvalues
    # Normalize gradient
    self.dinputs = self.dinputs / samples

class Activation_Softmax_Loss_CategoricalCrossentropy():
    # Creates activation and Loss function objects
    def __init__(self):
        self.activation = Activation_Softmax()
        self.loss = Loss_CategoricalCrossentropy()

    def forward(self, inputs, y_true):
        # Output layer's activation function
        self.activation.forward(inputs)
        # Set the output
        self.output = self.activation.output
        # Calculate and return loss value
        return self.loss.calculate(self.output, y_true)

    def backward(self, dvalues, y_true):
        # Number of samples
        samples = len(dvalues)
        # If Labels are one-hot encoded,
        # turn them into discrete values
        if len(y_true.shape) == 2:
            y_true = np.argmax(y_true, axis=1)
        # Copy so we can safely modify
        self.dinputs = dvalues.copy()
        # Calculate gradient
        self.dinputs[range(samples), y_true] -= 1
        # Normalize gradient
        self.dinputs = self.dinputs / samples

```

```

In [10]: class Optimizer_SGD:
    """
    Vanilla option - SGD
    Momentum option - if specified else default 0
    """
    def __init__(self, learning_rate=1.0, decay=0, momentum=0.):
        self.learning_rate = learning_rate
        self.current_learning_rate = learning_rate
        self.decay = decay
        self.iterations = 0
        self.momentum = momentum

    # Call once before any parameter updates

```

```

def pre_update_params(self):
    if self.decay:
        self.current_learning_rate = self.learning_rate * (1. / (1. + self.decay))

    # Update parameters
def update_params(self, layer):
    # If we use momentum
    if self.momentum:

        # If layer does not contain momentum arrays, create them
        # filled with zeros
        if not hasattr(layer, 'weight_momentums'):
            layer.weight_momentums = np.zeros_like(layer.weights)
            # If there is no momentum array for weights
            # The array doesn't exist for biases yet either.
            layer.bias_momentums = np.zeros_like(layer.biases)

        # Build weight updates with momentum - take previous
        # updates multiplied by retain factor and update with
        # current gradients
        weight_updates = self.momentum * layer.weight_momentums - self.current_learning_rate * layer.dweights
        layer.weight_momentums = weight_updates

        # Build bias updates
        bias_updates = self.momentum * layer.bias_momentums - self.current_learning_rate * layer.dbiases
        layer.bias_momentums = bias_updates

    # Vanilla SGD updates (as before momentum update)
    else:
        weight_updates = -self.current_learning_rate * layer.dweights
        bias_updates = -self.current_learning_rate * layer.dbiases

    # Update weights and biases using either
    # vanilla or momentum updates
    layer.weights += weight_updates
    layer.biases += bias_updates

# Call once after any parameter updates
def post_update_params(self):
    self.iterations += 1

```

In [11]:

```

dense1 = Layer_Dense(X.shape[1], 128, weight_regularizer_l2=5e-4, bias_regularizer_l2=5e-4)
activation1 = Activation_ReLU()
dense2 = Layer_Dense(128, 128)
activation2 = Activation_ReLU()
dense3 = Layer_Dense(128, 10)
loss_activation = Activation_Softmax_Loss_CategoricalCrossentropy()

optimizer = Optimizer_SGD(decay=0.01, momentum=0.5)

accuracies = []
losses = []
learning_rate = []

```

In [12]:

```

for epoch in range(1100):
    dense1.forward(X)
    activation1.forward(dense1.output)
    dense2.forward(activation1.output)

```

```
activation2.forward(dense2.output)
dense3.forward(activation2.output)
data_loss = loss_activation.forward(dense3.output, y)

regularization_loss = loss_activation.loss.regularization_loss(dense1) + loss_activation.loss.regularization_loss(dense3)

loss = data_loss + regularization_loss

predictions = np.argmax(loss_activation.output, axis=1)
if len(y.shape) == 2:
    y = np.argmax(y, axis=1)
accuracy = np.mean(predictions==y)

if not epoch % 10:
    accuracies.append(accuracy)
    losses.append(loss)
    learning_rate.append(optimizer.current_learning_rate)
    print(f'epoch: {epoch}, ' + f'acc: {accuracy:.3f}, ' + f'loss: {loss:.3f}, ' +
        f'Rr: {regularization_loss}')

loss_activation.backward(loss_activation.output, y)
dense3.backward(loss_activation.dinputs)
activation2.backward(dense3.dinputs)
dense2.backward(activation2.dinputs)
activation1.backward(dense2.dinputs)
dense1.backward(activation1.dinputs)

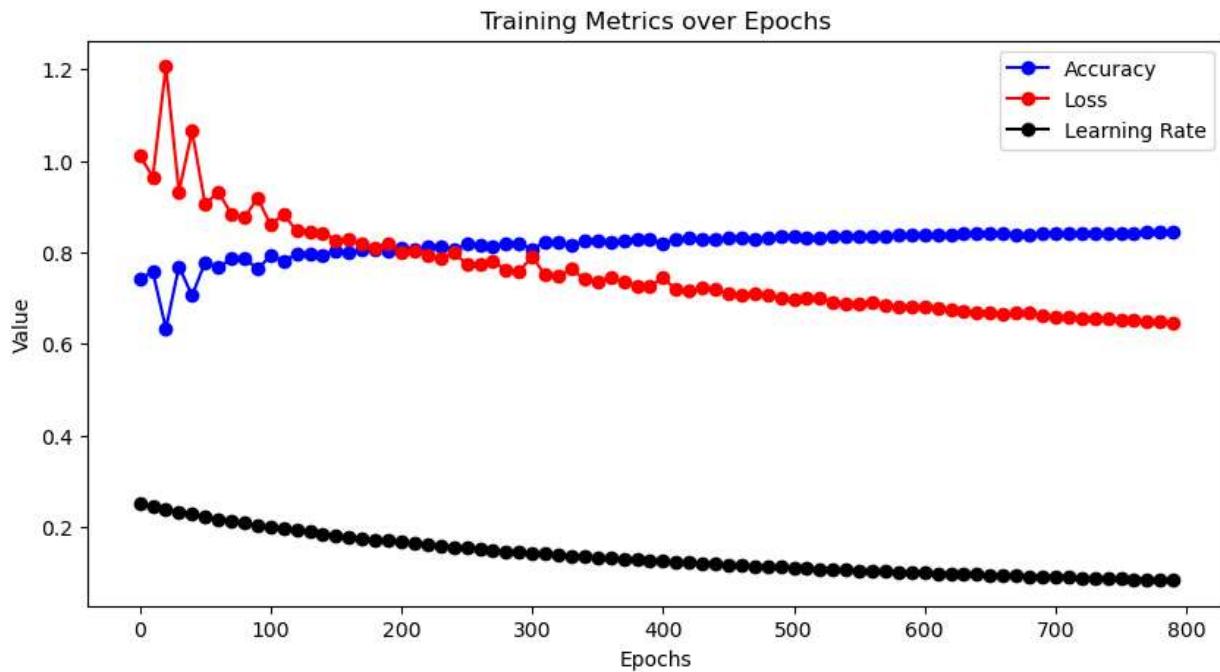
optimizer.pre_update_params()
optimizer.update_params(dense1)
optimizer.update_params(dense2)
optimizer.update_params(dense3)
optimizer.post_update_params()
```

epoch: 0, acc: 0.112, loss: 2.308, lr: 1.0Rr: 0.00503451968865524
epoch: 10, acc: 0.100, loss: 10.680, lr: 0.9174311926605504Rr: 0.00838227200825124
epoch: 20, acc: 0.238, loss: 2.456, lr: 0.8403361344537815Rr: 0.4002204596130479
epoch: 30, acc: 0.335, loss: 2.184, lr: 0.7751937984496123Rr: 0.406407055266603
epoch: 40, acc: 0.313, loss: 2.323, lr: 0.7194244604316546Rr: 0.41817819724849087
epoch: 50, acc: 0.185, loss: 2.690, lr: 0.6711409395973155Rr: 0.42146232069142076
epoch: 60, acc: 0.169, loss: 2.628, lr: 0.628930817610063Rr: 0.4136935871644017
epoch: 70, acc: 0.298, loss: 2.170, lr: 0.591715976331361Rr: 0.4059028981217895
epoch: 80, acc: 0.341, loss: 2.039, lr: 0.5586592178770949Rr: 0.3975730039535512
epoch: 90, acc: 0.203, loss: 2.439, lr: 0.5291005291005291Rr: 0.3899865106487447
epoch: 100, acc: 0.174, loss: 2.643, lr: 0.5025125628140703Rr: 0.3868441508671544
epoch: 110, acc: 0.216, loss: 2.289, lr: 0.47846889952153115Rr: 0.3797515158791232
epoch: 120, acc: 0.213, loss: 2.312, lr: 0.4566210045662101Rr: 0.372698010463934
epoch: 130, acc: 0.369, loss: 1.812, lr: 0.4366812227074236Rr: 0.3661404638099571
epoch: 140, acc: 0.313, loss: 2.024, lr: 0.41841004184100417Rr: 0.36015653418287336
epoch: 150, acc: 0.273, loss: 2.745, lr: 0.4016064257028112Rr: 0.354834598073053
epoch: 160, acc: 0.268, loss: 2.216, lr: 0.3861003861003861Rr: 0.363729285424908
epoch: 170, acc: 0.285, loss: 2.277, lr: 0.3717472118959108Rr: 0.3602163278317301
epoch: 180, acc: 0.376, loss: 1.781, lr: 0.35842293906810035Rr: 0.3555415370514555
epoch: 190, acc: 0.424, loss: 1.571, lr: 0.3460207612456747Rr: 0.350705861404834
epoch: 200, acc: 0.487, loss: 1.648, lr: 0.33444816053511706Rr: 0.3457487538790493
epoch: 210, acc: 0.293, loss: 2.296, lr: 0.3236245954692557Rr: 0.3413210042822297
epoch: 220, acc: 0.321, loss: 1.915, lr: 0.31347962382445144Rr: 0.3377538457315041
epoch: 230, acc: 0.559, loss: 1.449, lr: 0.303951367781155Rr: 0.33392780742332573
epoch: 240, acc: 0.524, loss: 1.388, lr: 0.2949852507374631Rr: 0.33096544813019124
epoch: 250, acc: 0.586, loss: 1.239, lr: 0.28653295128939826Rr: 0.3271496283947017
epoch: 260, acc: 0.684, loss: 1.153, lr: 0.2785515320334262Rr: 0.3235672674391777
epoch: 270, acc: 0.507, loss: 1.536, lr: 0.2710027100271003Rr: 0.3202406986157577
epoch: 280, acc: 0.715, loss: 1.088, lr: 0.2638522427440633Rr: 0.3174255490015097
epoch: 290, acc: 0.666, loss: 1.139, lr: 0.2570694087403599Rr: 0.3141771709173613
epoch: 300, acc: 0.742, loss: 1.012, lr: 0.2506265664160401Rr: 0.3112368047446666
epoch: 310, acc: 0.758, loss: 0.965, lr: 0.24449877750611249Rr: 0.30831208082879535
epoch: 320, acc: 0.634, loss: 1.206, lr: 0.23866348448687355Rr: 0.30548624722338996
epoch: 330, acc: 0.769, loss: 0.932, lr: 0.2331002331002331Rr: 0.30288957612238476
epoch: 340, acc: 0.705, loss: 1.065, lr: 0.2277904328018223Rr: 0.3002569911403449
epoch: 350, acc: 0.779, loss: 0.905, lr: 0.22271714922048996Rr: 0.2978217357830459
epoch: 360, acc: 0.766, loss: 0.932, lr: 0.2178649237472767Rr: 0.2953227540733903
epoch: 370, acc: 0.786, loss: 0.884, lr: 0.21321961620469085Rr: 0.29296502410304565
epoch: 380, acc: 0.788, loss: 0.875, lr: 0.20876826722338204Rr: 0.29057677395294296
epoch: 390, acc: 0.763, loss: 0.919, lr: 0.20449897750511245Rr: 0.28834437103459526
epoch: 400, acc: 0.793, loss: 0.860, lr: 0.2004008016032064Rr: 0.28615070258906034
epoch: 410, acc: 0.781, loss: 0.883, lr: 0.19646365422396858Rr: 0.2839843248900268
epoch: 420, acc: 0.796, loss: 0.847, lr: 0.19267822736030826Rr: 0.28190963366499644
epoch: 430, acc: 0.795, loss: 0.846, lr: 0.1890359168241966Rr: 0.27983635092609416
epoch: 440, acc: 0.795, loss: 0.843, lr: 0.1855287569573284Rr: 0.2778848358064881
epoch: 450, acc: 0.803, loss: 0.827, lr: 0.18214936247723132Rr: 0.27591801748527817
epoch: 460, acc: 0.801, loss: 0.828, lr: 0.17889087656529518Rr: 0.273988602461581
epoch: 470, acc: 0.805, loss: 0.819, lr: 0.17574692442882248Rr: 0.27220016713795486
epoch: 480, acc: 0.808, loss: 0.811, lr: 0.17271157167530224Rr: 0.2703865720905306
epoch: 490, acc: 0.802, loss: 0.820, lr: 0.16977928692699493Rr: 0.26861669483552764
epoch: 500, acc: 0.810, loss: 0.800, lr: 0.1669449081803005Rr: 0.26690238550154544
epoch: 510, acc: 0.808, loss: 0.803, lr: 0.16420361247947454Rr: 0.2652118463588427
epoch: 520, acc: 0.811, loss: 0.793, lr: 0.16155088852988692Rr: 0.2635828164462515
epoch: 530, acc: 0.814, loss: 0.786, lr: 0.1589825119236884Rr: 0.2619655119386929
epoch: 540, acc: 0.806, loss: 0.799, lr: 0.1564945226917058Rr: 0.26040678962867503
epoch: 550, acc: 0.818, loss: 0.774, lr: 0.15408320493066255Rr: 0.25886856412269804
epoch: 560, acc: 0.816, loss: 0.774, lr: 0.15174506828528073Rr: 0.25735475951600567
epoch: 570, acc: 0.812, loss: 0.781, lr: 0.14947683109118085Rr: 0.2559052643244058
epoch: 580, acc: 0.820, loss: 0.761, lr: 0.14727540500736377Rr: 0.25445608755202737
epoch: 590, acc: 0.820, loss: 0.759, lr: 0.14513788098693758Rr: 0.25303260182852055

epoch: 600, acc: 0.806, loss: 0.790, lr: 0.14306151645207438Rr: 0.2516672522069906
 epoch: 610, acc: 0.823, loss: 0.750, lr: 0.14104372355430184Rr: 0.2503171806493208
 epoch: 620, acc: 0.823, loss: 0.748, lr: 0.13908205841446453Rr: 0.2489739183251788
 epoch: 630, acc: 0.815, loss: 0.764, lr: 0.13717421124828533Rr: 0.24766286804043333
 epoch: 640, acc: 0.825, loss: 0.741, lr: 0.13531799729364005Rr: 0.2464009886874537
 epoch: 650, acc: 0.826, loss: 0.737, lr: 0.13351134846461948Rr: 0.2451360995633058
 epoch: 660, acc: 0.821, loss: 0.746, lr: 0.13175230566534915Rr: 0.24389374011743867
 epoch: 670, acc: 0.825, loss: 0.735, lr: 0.13003901170351104Rr: 0.24270128365512475
 epoch: 680, acc: 0.829, loss: 0.727, lr: 0.12836970474967907Rr: 0.24150960021136902
 epoch: 690, acc: 0.827, loss: 0.727, lr: 0.12674271229404308Rr: 0.2403316435284347
 epoch: 700, acc: 0.819, loss: 0.744, lr: 0.1251564455569462Rr: 0.23918524891273077
 epoch: 710, acc: 0.830, loss: 0.719, lr: 0.12360939431396786Rr: 0.2380717623508691
 epoch: 720, acc: 0.831, loss: 0.716, lr: 0.12210012210012208Rr: 0.23695576625227158
 epoch: 730, acc: 0.827, loss: 0.722, lr: 0.12062726176115804Rr: 0.23585843094784978
 epoch: 740, acc: 0.828, loss: 0.720, lr: 0.11918951132300357Rr: 0.23479644469833075
 epoch: 750, acc: 0.832, loss: 0.708, lr: 0.11778563015312131Rr: 0.23374401009231477
 epoch: 760, acc: 0.832, loss: 0.707, lr: 0.11641443538998836Rr: 0.23270052992162304
 epoch: 770, acc: 0.830, loss: 0.711, lr: 0.1150747986191024Rr: 0.23167483016860638
 epoch: 780, acc: 0.831, loss: 0.707, lr: 0.11376564277588169Rr: 0.2306769069114837
 epoch: 790, acc: 0.834, loss: 0.700, lr: 0.11248593925759279Rr: 0.22968870911631437
 epoch: 800, acc: 0.834, loss: 0.698, lr: 0.11123470522803114Rr: 0.2287106643692249
 epoch: 810, acc: 0.832, loss: 0.699, lr: 0.11001100110011001Rr: 0.22774802183629586
 epoch: 820, acc: 0.832, loss: 0.699, lr: 0.1088139281828074Rr: 0.22680670862269828
 epoch: 830, acc: 0.835, loss: 0.691, lr: 0.1076426264800861Rr: 0.22587869442193545
 epoch: 840, acc: 0.836, loss: 0.689, lr: 0.10649627263045792Rr: 0.22495956771052214
 epoch: 850, acc: 0.836, loss: 0.689, lr: 0.1053740779768177Rr: 0.22405349037474043
 epoch: 860, acc: 0.835, loss: 0.690, lr: 0.10427528675703858Rr: 0.2231637868140928
 epoch: 870, acc: 0.836, loss: 0.684, lr: 0.10319917440660475Rr: 0.22228897307046677
 epoch: 880, acc: 0.837, loss: 0.681, lr: 0.10214504596527067Rr: 0.22142303956061293
 epoch: 890, acc: 0.837, loss: 0.680, lr: 0.10111223458038422Rr: 0.22056818243880688
 epoch: 900, acc: 0.837, loss: 0.680, lr: 0.10010010010010009Rr: 0.21972637330096356
 epoch: 910, acc: 0.837, loss: 0.677, lr: 0.09910802775024777Rr: 0.21889786268828346
 epoch: 920, acc: 0.839, loss: 0.674, lr: 0.09813542688910697Rr: 0.21808011934389276
 epoch: 930, acc: 0.840, loss: 0.671, lr: 0.09718172983479105Rr: 0.2172717483608713
 epoch: 940, acc: 0.841, loss: 0.669, lr: 0.09624639076034648Rr: 0.2164731677002239
 epoch: 950, acc: 0.840, loss: 0.668, lr: 0.09532888465204957Rr: 0.21568463411760205
 epoch: 960, acc: 0.840, loss: 0.666, lr: 0.09442870632672333Rr: 0.2149063782384027
 epoch: 970, acc: 0.839, loss: 0.668, lr: 0.09354536950420955Rr: 0.21413904964009467
 epoch: 980, acc: 0.838, loss: 0.669, lr: 0.09267840593141798Rr: 0.21338735583826487
 epoch: 990, acc: 0.841, loss: 0.662, lr: 0.09182736455463728Rr: 0.21264512590920423
 epoch: 1000, acc: 0.842, loss: 0.660, lr: 0.09099181073703366Rr: 0.21190822497789233
 epoch: 1010, acc: 0.842, loss: 0.658, lr: 0.09017132551848513Rr: 0.21117964117771595
 epoch: 1020, acc: 0.842, loss: 0.657, lr: 0.08936550491510277Rr: 0.21046016692546368
 epoch: 1030, acc: 0.842, loss: 0.656, lr: 0.08857395925597873Rr: 0.20974990373548155
 epoch: 1040, acc: 0.841, loss: 0.655, lr: 0.08779631255487269Rr: 0.20904912304126
 epoch: 1050, acc: 0.842, loss: 0.653, lr: 0.08703220191470844Rr: 0.20835730436982514
 epoch: 1060, acc: 0.843, loss: 0.651, lr: 0.08628127696289906Rr: 0.20767388518449204
 epoch: 1070, acc: 0.843, loss: 0.650, lr: 0.0855431993156544Rr: 0.20699831315887868
 epoch: 1080, acc: 0.844, loss: 0.648, lr: 0.08481764206955046Rr: 0.2063304252037452
 epoch: 1090, acc: 0.844, loss: 0.646, lr: 0.08410428931875526Rr: 0.20567010121676432

```
In [17]: epochs = range(0, 800, 10)
plt.figure(figsize=(10, 5))
plt.plot(epochs, accuracies[30:], label='Accuracy', marker='o', color="blue")
plt.plot(epochs, losses[30:], label='Loss', marker='o', color='red')
plt.plot(epochs, learning_rate[30:], label='Learning Rate', marker='o', color='black')
plt.title('Training Metrics over Epochs')
plt.xlabel('Epochs')
plt.ylabel('Value')
plt.legend()
```

```
plt.savefig('training_metrics_plot.png')
plt.show()
```



In [14]: # Validate the model

```
dense1.forward(X_test)
activation1.forward(dense1.output)
dense2.forward(activation1.output)
activation2.forward(dense2.output)
dense3.forward(activation2.output)
data_loss = loss_activation.forward(dense3.output, y_test)

predictions = np.argmax(loss_activation.output, axis=1)
if len(y_test.shape) == 2:
    y_test = np.argmax(y, axis=1)
accuracy = np.mean(predictions==y_test)
print(f'validation, acc: {accuracy:.3f}, loss: {loss:.3f}')

validation, acc: 0.825, loss: 0.645
```

In []: